

GPU Occupancy Analysis with Gem5

Objectives

- Edit GPU occupancy analysis script from Module 6 to search for the correct stat names in gem5 stats files
- Adjust occupancy formulas to match with occupancy definitions
- Sweep through HIP parameters using HIP-samples code
- Compare GPU occupancy results between runs

Software Terminology

Nvidia/CUDA Terminology	AMD Terminology	Description
Streaming Multiprocessor	Compute Unit (CU)	One of many parallel vector processors in a GPU that contain parallel ALUs. All waves in a workgroups are assigned to the same CU.
Kernel	Kernel	Functions launched to the GPU that are executed by multiple parallel workers on the GPU. Kernels can work in parallel with CPU.
Warp	Wavefront	Collection of operations that execute in lockstep, run the same instructions, and follow the same control-flow path. Individual lanes can be masked off. Think of this as a vector thread. A 64-wide wavefront is a 64-wide vector op.
Thread Block	Workgroup	Group of wavefronts that are on the GPU at the same time. Can synchronize together and communicate through local memory.
Thread	Work Item / Thread	Individual lane in a wavefront. On AMD GPUs, must run in lockstep with other work items in the wavefront. Lanes can be individually masked off. GPU programming models can treat this as a separate thread of execution, though you do not necessarily get forward sub-wavefront progress.

Original GPU_occupancy_analyzer.py Issues

```
def _get_active_blocks(self):
    """Extract active block count"""
    for stat_name, stat_data in self.stats.items():
        if 'block' in stat_name.lower() and 'active' in stat_name.lower():
            return stat_data['value']

    # Look for work-group stats (AMD terminology)
    for stat_name, stat_data in self.stats.items():
        if 'work_group' in stat_name.lower() or 'workgroup' in stat_name.lower():
            return stat_data['value']

    return None

def _get_total_cycles(self):
    """Get total simulation cycles"""
    for stat_name, stat_data in self.stats.items():
        if 'sim_ticks' in stat_name.lower() or 'total_cycles' in stat_name.lower():
            return stat_data['value']
    return None

def _get_active_cycles(self):
    """Get active execution cycles"""
    for stat_name, stat_data in self.stats.items():
        if 'active_cycles' in stat_name.lower() or 'busy_cycles' in stat_name.lower():
            return stat_data['value']
    return None

def get_busy_cycles(self):
```

- Matching stat names did not exist, or did exist but did not represent what we wanted to look for

Changes to GPU_config

```
99      }
00  parser.add_argument(
01      "--simds-per-cu", type=int, default=4, help="SIMD units per CU"
02  )
03  parser.add_argument(
04      "--wfs-per-simd",
05      type=int,
06      default=10,
07      help="Number of WF slots per SIMD",
08  )
09  parser.add_argument(
10      "-u",
11      "--num-compute-units",
12      type=int,
13      default=4,
14      help="number of GPU compute units",
15  ),
16  # issue period per SIMD unit: number of cycles before issuing another vector
17  parser.add_argument(
18      "--issue-period",
19      type=int,
20      default=4,
21      help="Number of cycles per vector instruction issue period",
22  )
```

- wax_warps_per_sm : 64 -> 40
- Warp_size: 64
- Max_threads_per_sm: 2048 -> 2560
- Num_sms: 64 -> 4
- Max IPC: [1 / (instruction issue period)] times SIMDs per cu = 1 instr/cycle
- Max Blocks per CU: 40

Hardware feature support	RDNA1
Maximum number of resident blocks per compute unit	40 [1]
Maximum number of resident wavefronts per compute unit	40 [1]

GPU Config

```
# GPU architecture parameters
self.gpu_config = {
    'max_warps_per_sm': 40,          # Maximum warps per SM (typical for modern GPUs)
    'warp_size': 64,                # AMD waveform size (64) vs NVIDIA warp size (32)
    'max_threads_per_sm': 2560,      # Maximum threads per SM
    'num_sms': 4,                  # Number of streaming multiprocessors/compute units #64
    'max_blocks_per_sm': 40,         # Maximum blocks per SM. changed from 32
    'shared_memory_per_sm': 65536,   # Shared memory per SM in bytes
    'registers_per_sm': 8192,        # Number of registers per SM
    'peak_mem_bandwidth_gbps': 50,   #default used to be 900 GB/s, but that is for modern
}
```

Thread-based Occupancy

system.cpu3.CUs1.vALUInsts	264352	# Number of vector ALU insts issued. (Unspecified)
system.cpu3.CUs1.vALUInstsPerWF	16	# The avg. number of vector ALU insts issued per-wavefront. (Unspecified)
system.cpu3.CUs1.sALUInsts	66088	# Number of scalar ALU insts issued. (Unspecified)
system.cpu3.CUs1.sALUInstsPerWF	4	# The avg. number of scalar ALU insts issued per-wavefront. (Unspecified)
system.cpu3.CUs1.instCyclesVALU	264352	# Number of cycles needed to execute VALU insts. (Unspecified)
system.cpu3.CUs1.instCyclesSALU	66088	# Number of cycles needed to execute SALU insts. (Unspecified)
system.cpu3.CUs1.threadCyclesVALU	4229632	# Number of thread cycles used to execute vector ALU ops. Similar to instCyclesVALU.
system.cpu3.CUs1.vALUUtilization	25	# Percentage of active vector ALU threads in a wave. (Unspecified)
system.cpu3.CUs1.vALUUtilization	0	# Number of L1 cache misses in a wave. (Unspecified)
system.cpu3.CUs1.numMairReadCasOps	0	# number of compare and swap operations that failed
system.cpu3.CUs1.completedWFs	16522	# number of completed wavefronts (Unspecified)
system.cpu3.CUs1.completedWGps	16522	# number of completed workgroups (Unspecified)
system.cpu3.CUs1.headTailLatency::bucket_size	10000	# ticks between first and last cache block access

- Thread occupancy = active threads per CU / max threads per CU
- Thread occupancy = [(%active_threads)*(wavefront_size)*(wavefronts_per_block)*(blocks_per_cu)] / (max_threads_per_cu)
- %active_threads = vALUUtilization value in stats file
- Wavefront_size = warp_size in gpu_config = 64
- Wavefronts_per_block = # completed wavefronts / # completed workgroups
- Blocks_per_cu = min(blocks_per_sm_threads, max_blocks_per_sm)
- Blocks_per_sm_threads = max_threads_per_sm / threads_per_block
- Threads_per_block = wavefronts_per_block * waveform_size
- Max_blocks_per_sm = parameter in gpu_config = 40
- Max_threads_per_sm = parameter in gpu_config = 2560

Block-based Occupancy

```
# Thread-based limit on active blocks per SM  
blocks_per_sm_threads = max_threads_per_sm // int(threads_per_block)  
  
blocks_per_sm = min(blocks_per_sm_threads, max_blocks_per_sm)  
  
# --- Theoretical block occupancy (% of max blocks per SM) ---  
block_occ = (blocks_per_sm / max_blocks_per_sm) * 100.0
```

- Block occupancy = blocks per CU / max blocks per CU
- Wavefront_size = warp_size in gpu_config = 64
- Wavefronts_per_block = # completed wavefronts / # completed workgroups
- Blocks_per_cu = min(blocks_per_sm_threads, max_blocks_per_sm)
- Blocks_per_sm_threads = max_threads_per_sm / threads_per_block
- Threads_per_block = wavefronts_per_block * wavefront_size
- Max_blocks_per_sm = parameter in gpu_config = 40
- Max_threads_per_sm = parameter in gpu_config = 2560

Theoretical Warp-based Occupancy

```
# --- Theoretical warp occupancy ---
#max_warp_per_sm_thread_limit = max_threads_per_sm // warp_size #just going to do integer division
warps_per_block = wavefronts_per_block
warps_per_sm = blocks_per_sm * warps_per_block

warp_occ = (warps_per_sm / self.gpu_config('max_warp_per_sm')) * 100.0
```

- Warp occupancy = active warps / max warps per CU
- Warps_per_sm = blocks_per_cu * wavefronts_per_block
- Wavefront_size = warp_size in gpu_config = 64
- Wavefronts_per_block = # completed wavefronts / # completed workgroups
- Blocks_per_cu = min(blocks_per_sm_threads, max_blocks_per_sm)
- Blocks_per_sm_threads = max_threads_per_sm / threads_per_block
- Threads_per_block = wavefronts_per_block * waveform_size
- Max_blocks_per_sm = parameter in gpu_config = 40
- Max_threads_per_sm = parameter in gpu_config = 2560

Achieved Warp-based Occupancy

```
system.cpu3.CUs0.waveLevelParallelism::samples      16306
system.cpu3.CUs0.waveLevelParallelism::mean        38.952165
system.cpu3.CUs0.waveLevelParallelism::stdev       1.121360
system.cpu3.CUs0.waveLevelParallelism::underflows    0     0.003
```

- Achieved Warp occupancy = achieved active warps / max warps per CU
- waveLevelParallelism::mean avg over all CUs = achieved active wavefronts

```
# wave level parallelism: count of active waves at wave launch (Unspecified)
# wave level parallelism: count of active waves at wave launch (Unspecified)
# wave level parallelism: count of active waves at wave launch (Unspecified)
# wave level parallelism: count of active waves at wave launch (Unspecified)

def _calculate_wave_occupancy(self):
    """Calculate achieved warp/wave occupancy from waveLevelParallelism stats."""
    metrics = {}

    max_waves_per_cu = self.gpu_config["max_warp_per_sm"] # 40 for gfx9
    num_cus = self.gpu_config["num_sms"] # 4 in gpu_config

    active_waves_means = []

    for cu_idx in range(num_cus):
        key = f"system.cpu3.CUs{cu_idx}.waveLevelParallelism::mean"
        entry = self.stats.get(key)
        if entry is None:
            continue

        mean_waves = entry["value"]
        if mean_waves <= 0:
            continue

        active_waves_means.append(mean_waves)

    if not active_waves_means:
        return metrics # nothing found

    # Average over CUs
    avg_active_waves = sum(active_waves_means) / len(active_waves_means)

    warp_occ = (avg_active_waves / max_waves_per_cu) * 100.0
    metrics["Achieved Warp Occupancy (%)"] = min(warp_occ, 100.0)

    return metrics["Achieved Warp Occupancy (%)"]
```

Cycle-based Occupancy / SM Utilization

system.cpu3.CUsl.ExecStage.numTransActiveIdle	263531	# number of CU transitions from active to idle (Unspecified)
system.cpu3.CUsl.ExecStage.numCyclesWithNoIssue	3549175	# number of cycles the CU issues nothing (Unspecified)
system.cpu3.CUsl.ExecStage.numCyclesWithInstrIssued	377537	# number of cycles the CU issued at least one instruction (Unspecified)
system.cpu3.CUsl.ExecStage.spc::samples	3926712	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::mean	0.117813	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::stdev	0.396581	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::underflows	0	0.00% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::0	3549175	90.39% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::1	310258	7.90% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::2	51307	1.31% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::3	14169	0.36% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::4	1778	0.05% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::5	25	0.00% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::6	0	0.00% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::7	0	0.00% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::8	0	0.00% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::overflows	0	0.00% # Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::min_value	0	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::max_value	5	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.spc::total	3926712	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::samples	263532	# duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::mean	13.467719	# duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::stdev	67.294193	# duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::underflows	0	0.00% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::0-4	229604	87.13% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::5-9	8203	3.11% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::10-14	1873	0.71% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::15-19	1414	0.54% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::20-24	1082	0.41% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::25-29	941	0.36% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::30-34	1330	0.50% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::35-39	1977	0.75% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::40-44	1786	0.68% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::45-49	1896	0.72% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::50-54	999	0.38% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::55-59	6187	2.35% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::60-64	105	0.04% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::65-69	133	0.05% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::70-74	154	0.06% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::overflows	5848	2.22% 100.00% # duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::min_value	1	# duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::max_value	3198	# duration of idle periods in cycles (Unspecified)
system.cpu3.CUsl.ExecStage.idleDur::total	263532	# duration of idle periods in cycles (Unspecified)

Cycle-based occupancy = cycles with instruction issued / total cycles

IPC-based Occupancy

```
5     # issue period per SIMD unit: number of cycles before issuing another vector
6     parser.add_argument(
7         "--issue-period",
8         type=int,
9         default=4,
10        help="Number of cycles per vector instruction issue period",
11    )
12    parser.add_argument(/
```

- Max IPC per CU = $(1 / \text{SIMD issue period}) * (\# \text{SIMDs} / \text{CU}) = (\frac{1}{4})(4) = 1 \text{ instr/cycle}$
- Achieved IPC = $\text{sum}[(\text{ipc} * \text{totalCycles}) \text{ for each CU}] / \text{sum}(\text{totalcycles all CUs})$

```
def _calculate_instruction_throughput(self):
    """Calculate instruction-based occupancy metrics using per-CU IPC"""
    metrics = {}

    cu_ipcs = []
    cu_cycles = []

    # assume 4 CUs: CUs0..CUs3
    for cu_idx in range(4):
        ipc_key = f"system.cpu3.CUs{cu_idx}.ipc"
        cycles_key = f"system.cpu3.CUs{cu_idx}.totalCycles"

        ipc_entry = self.stats.get(ipc_key)
        cycles_entry = self.stats.get(cycles_key)

        if not ipc_entry or not cycles_entry:
            continue

        ipc = ipc_entry["value"]
        cycles = cycles_entry["value"]

        # skip NaNs / zero cycles
        if cycles <= 0:
            continue
        if isinstance(ipc, float) and math.isnan(ipc):
            continue

        cu_ipcs.append(ipc)
        cu_cycles.append(cycles)

    if cu_cycles:
        # total instructions across all CUs
        total_insts = sum(ipc * cyc for ipc, cyc in zip(cu_ipcs, cu_cycles))
        print(f"total insts: {str(total_insts)}")
        # total CU-cycles across all CUs
        total_cu_cycles = sum(cu_cycles)
        print(f"total cu cycles: {str(total_cu_cycles)}")

        ipc_avg = total_insts / total_cu_cycles
        metrics["Instructions Per Cycle (avg per CU)"] = ipc_avg

        # IPC-based occupancy vs per-CU theoretical max
        theoretical_max_ipc_per_cu = 1.0 # assumed peak IPC per CU
        ipc_occupancy = (ipc_avg / theoretical_max_ipc_per_cu) * 100.0
        metrics["IPC-based Occupancy (%)"] = min(ipc_occupancy, 100.0)

    return metrics
```

```
system.cpu3.CUs0.numVecOpsExecutedInW00PF          0
system.cpu3.CUs0.totalCycles                         3926616
system.cpu3.CUs0.vpc                                 1.860403
system.cpu3.CUs0.vpc_f16                            0
system.cpu3.CUs0.vpc_f32                            0
system.cpu3.CUs0.vpc_f64                            0
system.cpu3.CUs0.ipc                                0.116275
```

```
# number of two op FP vec Ops executed (e.g. wr size/inst) (
# number of cycles the CU ran for (Unspecified)
# Vector Operations per cycle (this CU only) (Unspecified)
# F16 Vector Operations per cycle (this CU only) (Unspecified)
# F32 Vector Operations per cycle (this CU only) (Unspecified)
# F64 Vector Operations per cycle (this CU only) (Unspecified)
# Instructions per cycle (this CU only) (Unspecified)
```

Matrix Transpose with Varying Block Dimensions

(using gpu_occupancy_analyzer_v9.py)

MatrixTranspose

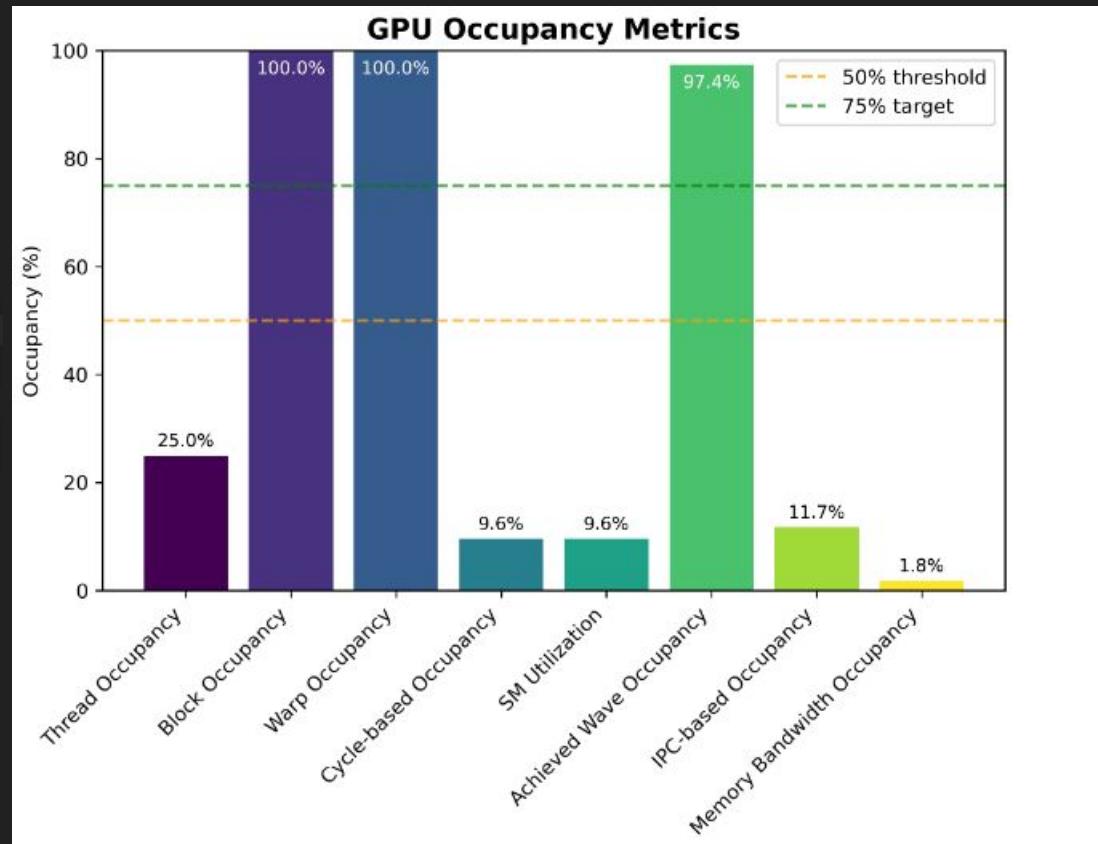
Width=1024

Threads per block X = 4

Threads per block Y = 4

Threads per block Z = 1

System.CPU.CU.S0.CompletedWPs	16306
System.CPU.CU.S0.CompletedWGs	16306
System.CPU.CU.S1.VALUUtilization	25
System.CPU.CU.S1.ThreadLevelTests	0
System.CPU.CU.S1.WaveLevelParallelism::mean	38.952790



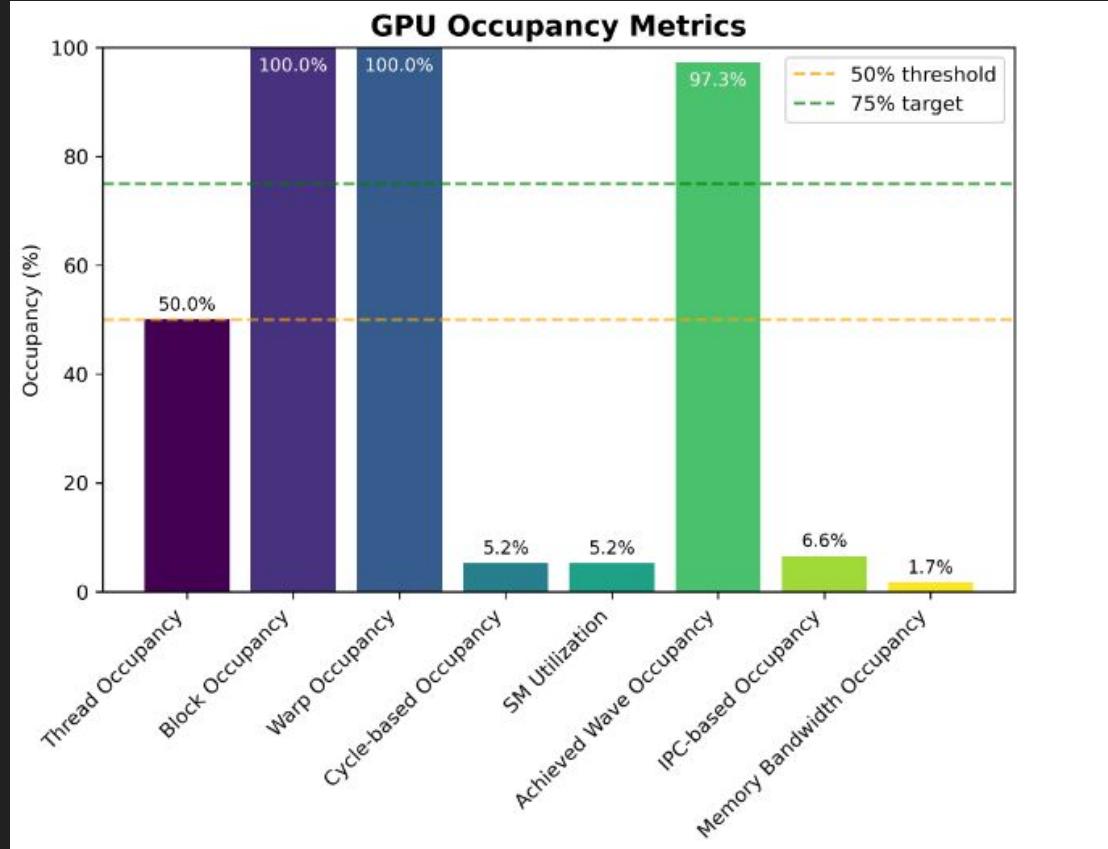
MatrixTranspose

Width=1024

Threads per block X = 8

Threads per block Y = 4

Threads per block Z = 1



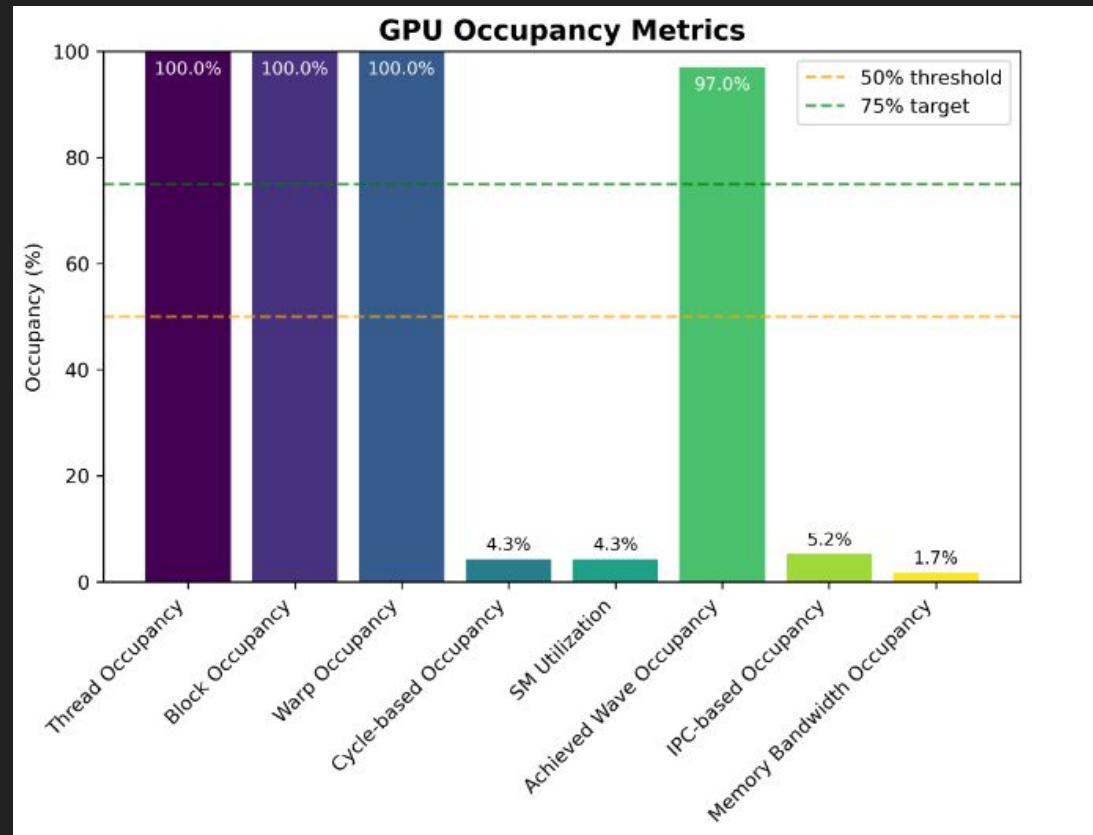
MatrixTranspose

Width=1024

Threads per block X = 8

Threads per block Y = 8

Threads per block Z = 1



MatrixTranspose

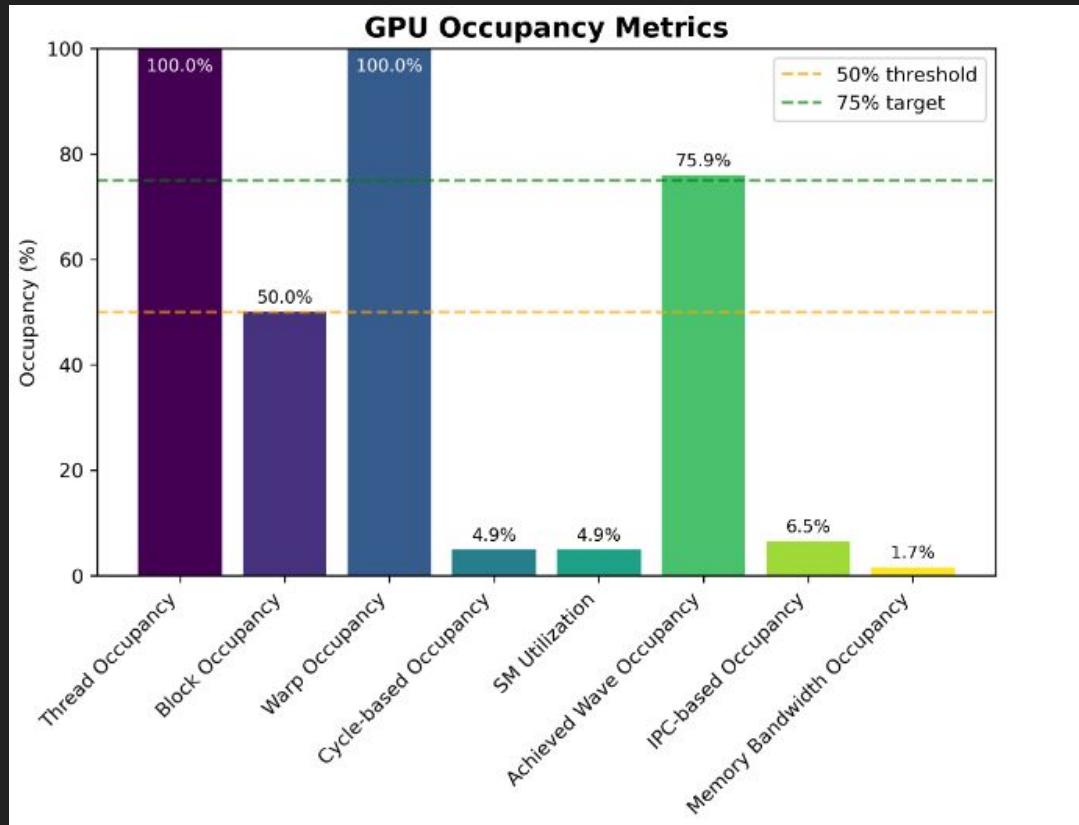
Width=1024

Threads per block X = 16

Threads per block Y = 8

Threads per block Z = 1

system.cpu3.CUs0.completedWfs	4200
system.cpu3.CUs0.completedWGs	2100



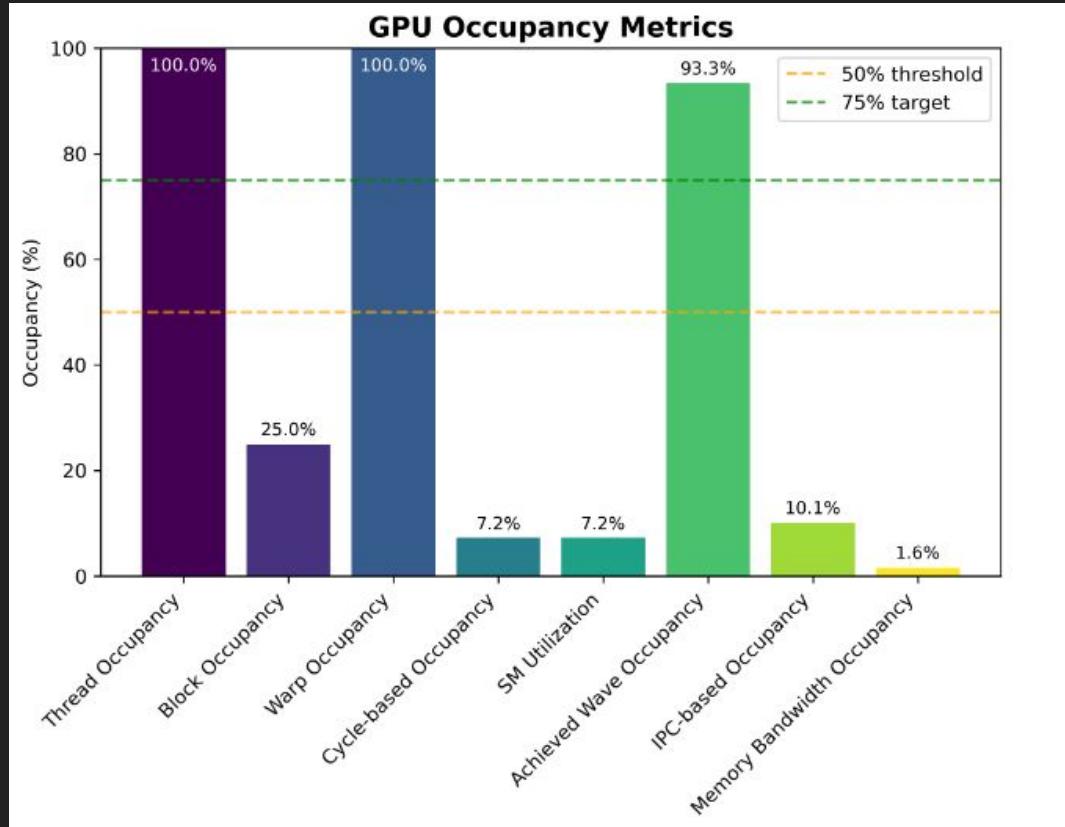
MatrixTranspose

Width=1024

Threads per block X = 16

Threads per block Y = 16

Threads per block Z = 1



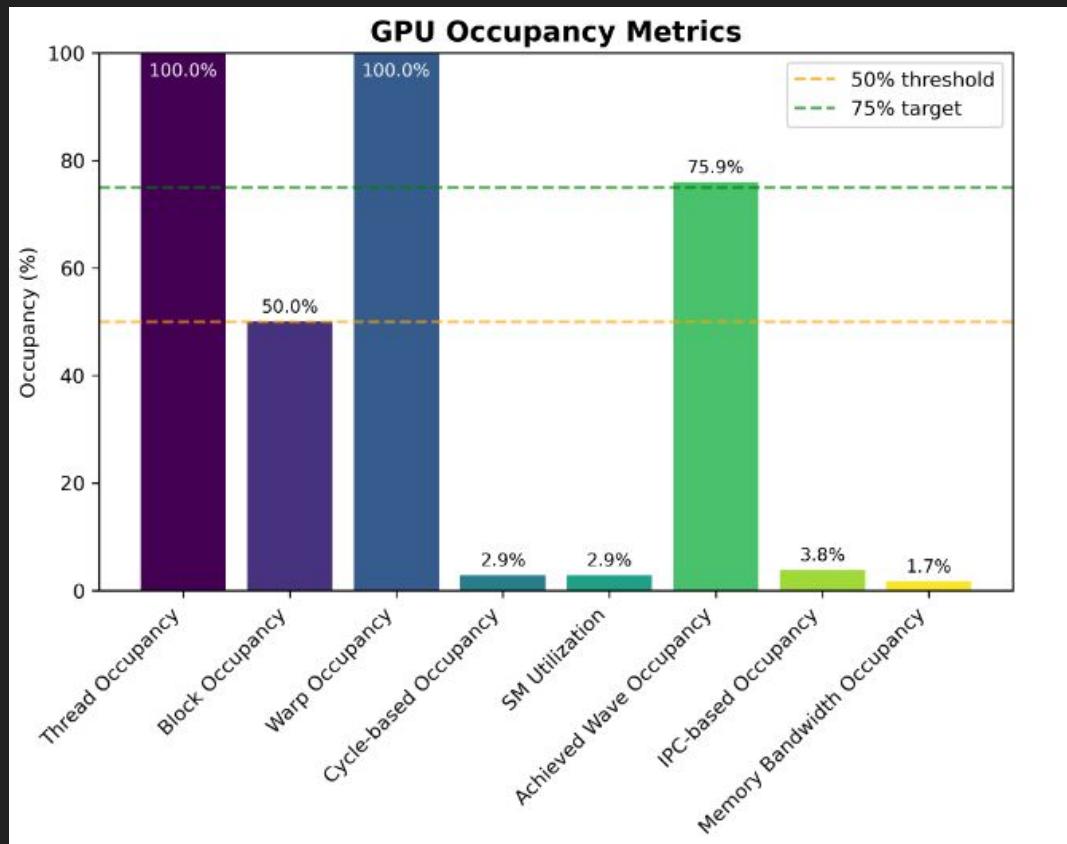
MatrixTranspose

Width=1024

Threads per block X = 32

Threads per block Y = 4

Threads per block Z = 1



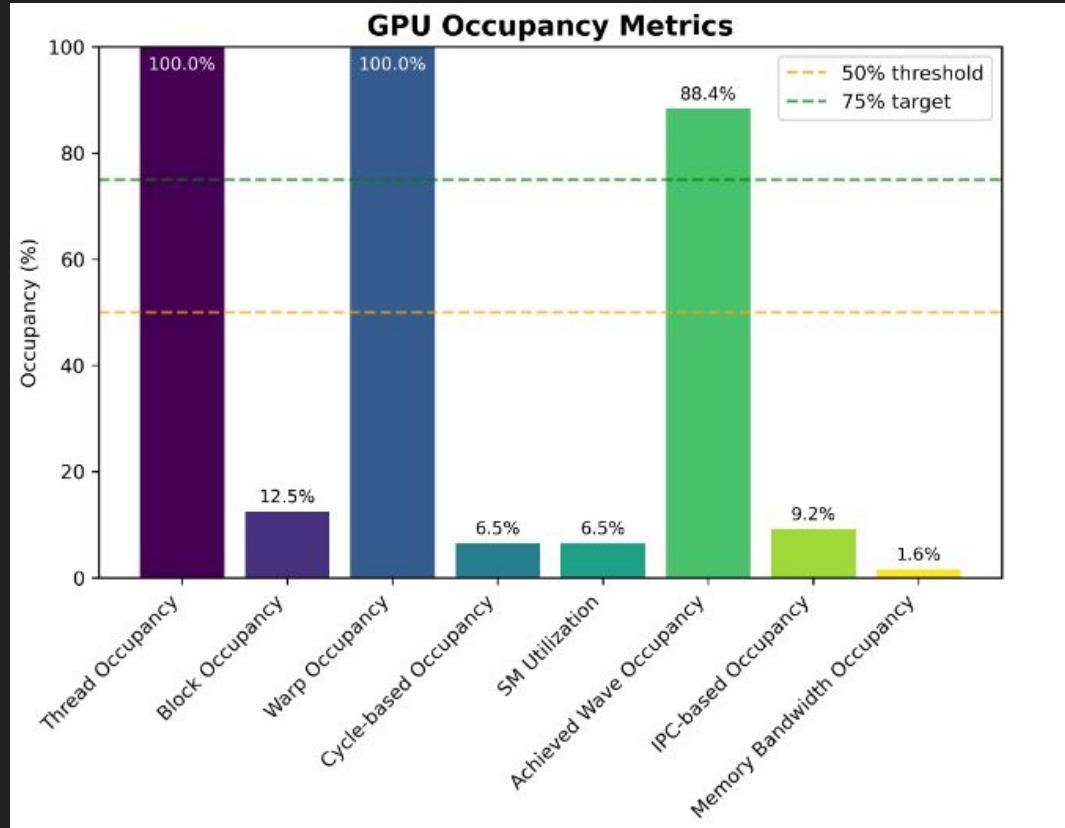
MatrixTranspose

Width=1024

Threads per block X = 32

Threads per block Y = 16

Threads per block Z = 1



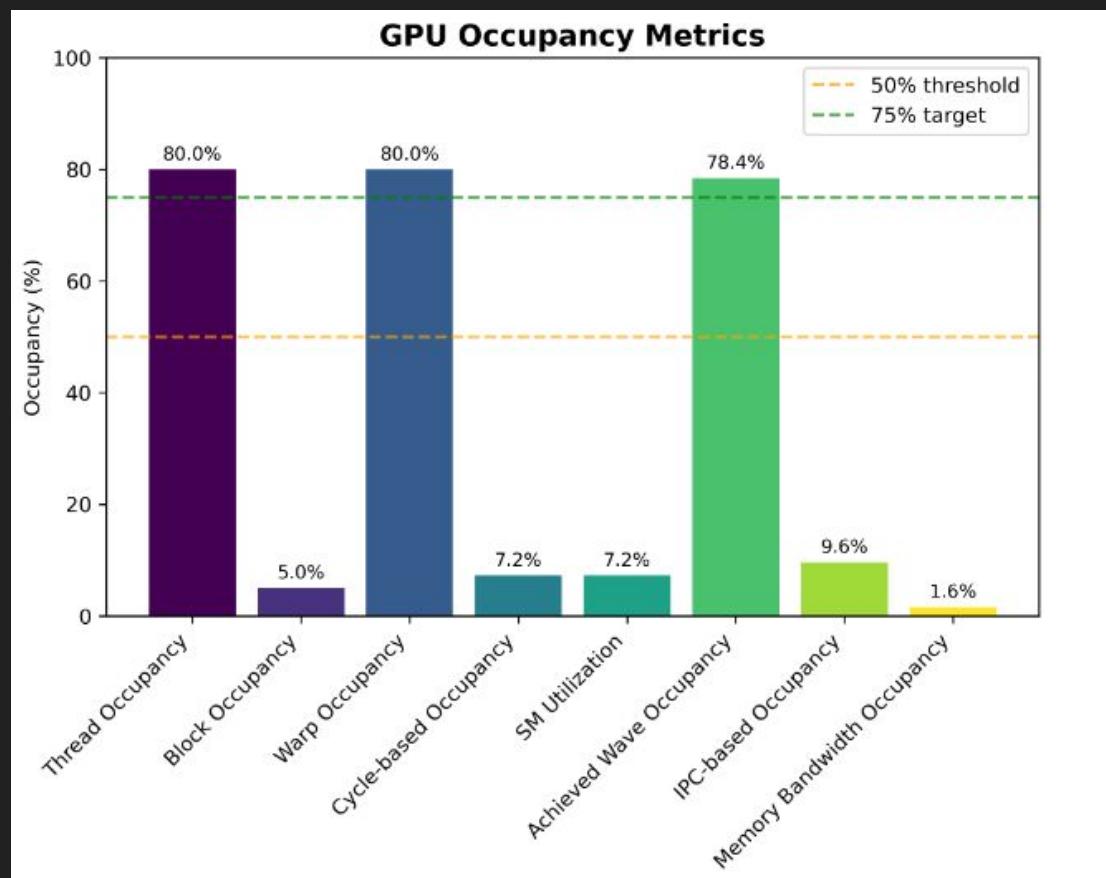
MatrixTranspose

Width=1024

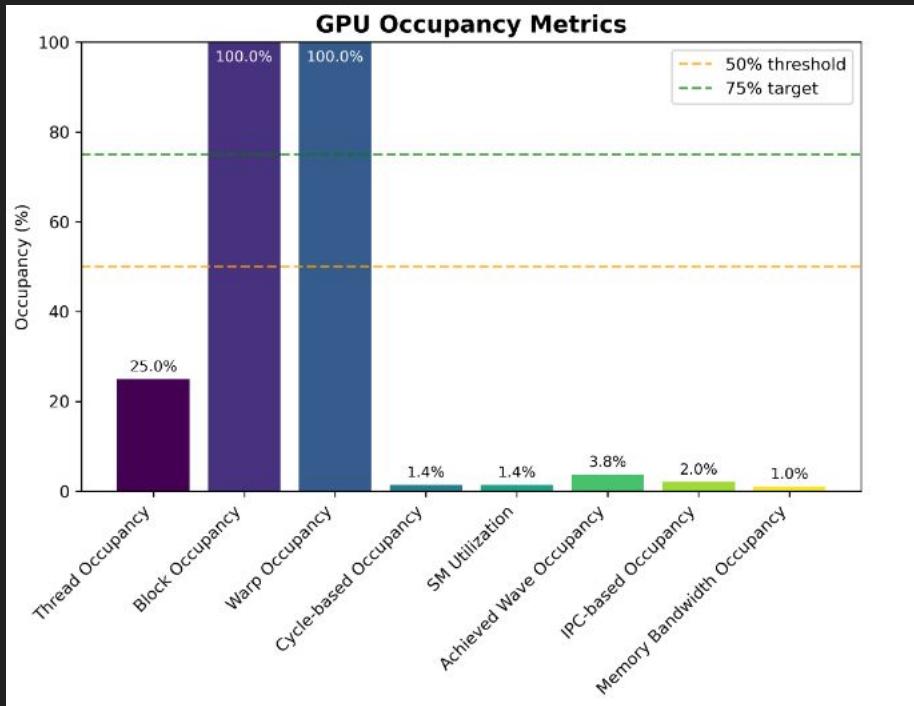
Threads per block X = 32

Threads per block Y = 32

Threads per block Z = 1

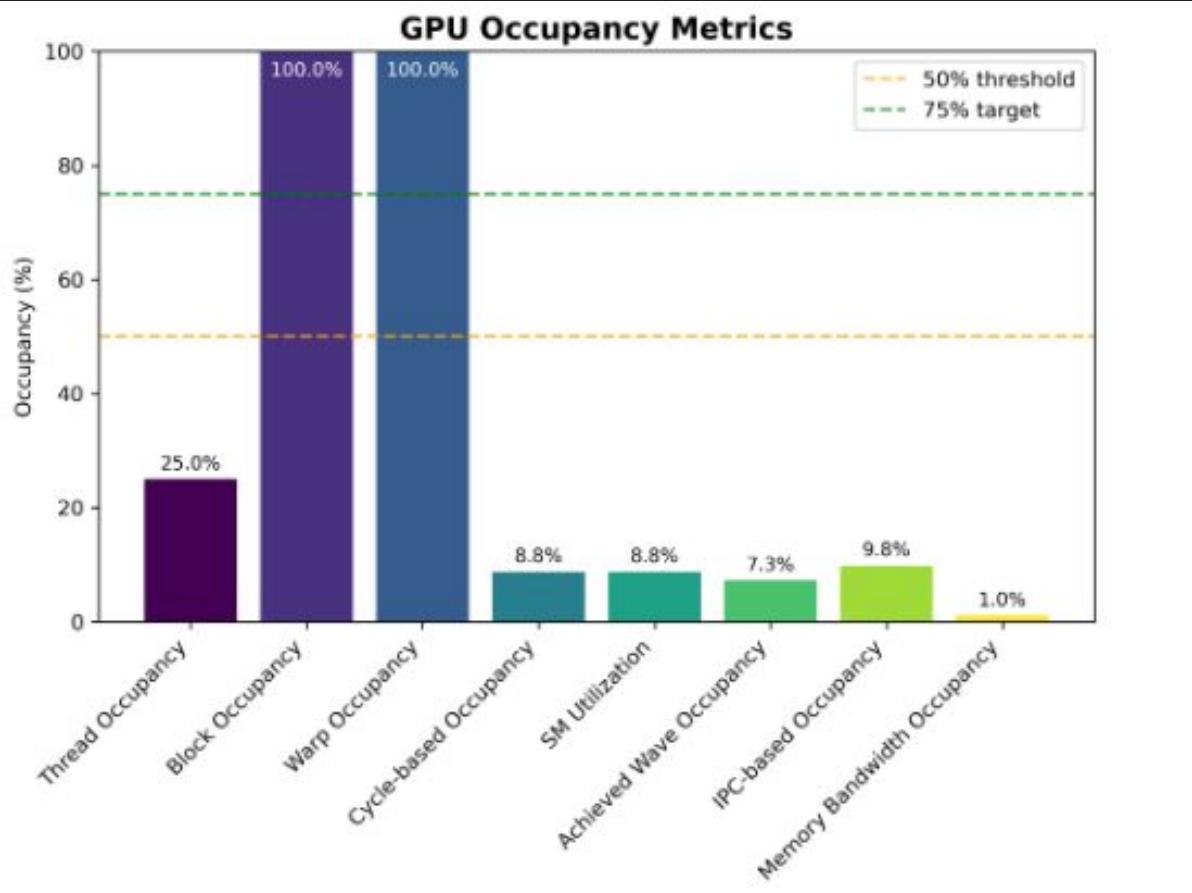


dynamic_shared



system.cpu3.CUs3.ExecStage.numCyclesWithNoIssue	6949	# number of cycles the CU issues nothing (Unspecified)
system.cpu3.CUs3.ExecStage.numCyclesWithInstrIssued	99	# number of cycles the CU issued at least one instruction (Unspecified)
system.cpu3.CUs3.ExecStage.spc::samples	7048	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUs3.ExecStage.spc::mean	0.019864	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUs3.ExecStage.spc::stdev	0.201181	# Execution units active per cycle (Exec unit=SIMD,MemPipe) (Unspecified)
system.cpu3.CUs3.numAlignedCASOps	0	# Number of compare and swap operations that fail
system.cpu3.CUs3.completedWfs	4	# number of completed wavefronts (Unspecified)
system.cpu3.CUs3.completedWGs	4	# number of completed workgroups (Unspecified)
system.cpu3.CUs3.headTailLatency::bucket_size	10000	# ticks between first and last cache block

sharedMemory



Conclusion

- Fixed most of the GPU occupancy calculations and stat parsing in the occupancy analyzer script
- Discussed occupancy formulas and important metrics from gem5 stats file
- Looked at how block size affects occupancy
- Resource utilization affecting occupancy
- Further work:
 - Assessing whether cycle based occupancy is calculated like it should be
 - Find the best metrics in the stats file for showing resource utilization and shared memory or LDS allocation and usage
 - Could help identify resource bottlenecks
 - Compare results of running on AMD hardware with the Radeon GPU profiler
- Github: https://github.com/alex-keist/EN525_712_gem5_gpu_occupancy

Sources

- https://www.olcf.ornl.gov/wp-content/uploads/2019/09/AMD_GPU_HIP_training_20190906.pdf
- <https://gpuopen.com/learn/occupancy-explained/>
- <https://rocm.docs.amd.com/en/latest/reference/gpu-arch-specs.html>
- <https://rocm.docs.amd.com/projects/HIP/en/latest/index.html>
- https://rocm.docs.amd.com/projects/HIP/en/latest/reference/hardware_features.html
- https://www.olcf.ornl.gov/wp-content/uploads/2019/10/ORNL_Application_Readiness_Workshop-AMD_GPU_Basics.pdf
- <https://github.com/gem5bootcamp/2024/blob/main/slides/04-GPU-model/gpu-slides.pdf>